

Scheduling Bag-of-Tasks in Clouds Using Spot and Burstable Virtual Machines

Luan Teylo^{ID}, Luciana Arantes, Pierre Sens^{ID}, and Lúcia Maria de A. Drummond^{ID}

Abstract—Cloud providers offer several types of Virtual Machines (VMs) in diverse markets, with different guarantees in terms of availability and reliability. Among them, the most popular market models are the on-demand and the spot. On-demand VMs are allocated for a fixed cost per time, and their availability is ensured during the whole execution. On the other hand, in the spot market, VMs are offered with a huge discount, but their availability fluctuates according to cloud's current demand that can terminate or hibernate a spot VM at any time. Furthermore, to cope with workload variations, cloud providers have also introduced the concept of burstable VMs, which can burst up their CPU performance during a limited period of time. In this work, we present the Burst Hibernation-Aware Dynamic Scheduler (Burst-HADS), a framework that executes Bag-of-Tasks applications with deadline constraints by exploiting both spot and on-demand burstable VMs, aiming at minimizing both the monetary cost and the execution time. Performance results on Amazon EC2 show that Burst-HADS reduces the monetary cost and meets the application deadline even in spot hibernation scenarios, when compared to other approaches from the related literature which uses only spot and non-burstable on-demand instances.

Index Terms—Cloud computing, BoT scheduling, burstable VMs, spot VM hibernation

1 INTRODUCTION

IN the past few years, cloud computing has emerged as an attractive solution to execute different applications. It brings several advantages compared with dedicated infrastructure, as, for example, a significant reduction in operational costs. However, in cloud environments, besides the usual goal of minimizing the application's execution time, it is also essential to reduce the monetary cost because even though in the cloud, computational resources are virtually infinite, the user's budget is not.

Cloud platforms enable users to dynamically acquire computational resources wrapped as Virtual Machines (VMs), that can be selected by the users according to their application requirements (CPU, memory, I/O, etc.) in a pay-as-you-use price model. Furthermore, cloud providers offer VMs in different contract models, with different guarantees in terms of availability and volatility. For instance, in Amazon EC2, there are three main contract models (also called markets): i) *reserved* market, where the user pays an upfront price, guaranteeing long-term availability; ii) *on-demand* market which is allocated for specific periods of time, and incurs a fixed cost per unit time of use, ensuring

the availability of the instance during this period; iii) *spot* market in which unused resources are available up to 90% discount when compared to the on-demand model, but such resources can be requested at any time.

In the three markets, there exist a wide range of VM types that suit different user requirements. According to Amazon Web Service (AWS), "Instance types comprise varying combinations of virtual CPUs (vCPUs), memory, storage, and networking capacity and give you the flexibility to choose the appropriate mix of resources for your applications. Each instance type includes one or more instance sizes, allowing you to scale your resources to the requirements of your target workload."¹ Instance types are grouped into families based on their respective use cases. For example, the compute-optimized instances (C3, C4, and C5) are ideal for compute-bound applications that require high-performance processors.

Regarding the spot market, the availability of its VMs fluctuates according to the cloud's current demand. If there are not enough resources to meet clients' requests, the cloud provider can interrupt a spot VM (temporarily or definitively). Despite the risk of unavailability, the main advantage of spot VMs is that their costs are much lower than on-demand VMs since the user requests unused instances at steep discounts. An interrupted spot VM instance can either terminate or hibernate. If the VM will be terminated, the cloud provider warns the user two minutes before its interruption. On the other hand, hibernated VM instances are frozen immediately after noticing the user. In this case, EC2 saves the VM instance memory and context in the root of EC2 Block Storage (EBS) volume, and during the VM's interruption period, the user is only charged for the EBS storage use. EC2 resumes the hibernated spot instance, reloading the saved memory and context, only when there

- Luan Teylo and Lúcia Maria de A. Drummond are with the Institute of Computing, Federal Fluminense University, Niterói, State of Rio de Janeiro 24220-900, Brazil. E-mail: luan-teylo@gmail.com, lucia@ic.uff.br.
- Luciana Arantes and Pierre Sens are with Sorbonne Université, CNRS, Inria, 75006 Paris, France. E-mail: {luciana.arantes, pierre.sens}@lip6.fr.

Manuscript received 19 February 2021; revised 14 October 2021; accepted 1 November 2021. Date of publication 8 November 2021; date of current version 8 March 2023.

This research was supported by Project No. CNPq/AWS 440014/2020-4, Brazil and by the Programa Institucional de Internacionalização (PrInt) from CAPES as part of the Project REMATCH under Project No. 88887.310261/2018-00.

(Corresponding author: Luan Teylo.)

Recommended for acceptance by A. Zomaya.

Digital Object Identifier no. 10.1109/TCC.2021.3125426

1. <https://aws.amazon.com/ec2/instance-types>

is enough available resource whose price is lower than the maximum one, with which the user agreed to be charged.

Besides the markets, some leading cloud providers (e.g., Microsoft Azure, Amazon EC2), introduced in the last years the concept of a burstable VM that can sprint its performance during a limited period of time to cope with sudden workload variations. By operating on a CPU credit regime that controls the processing power offered to users, burstable VM instances can use 100% of the VM's processing power (burst mode) or only a fraction of that processing power (baseline mode), depending on credits. They can then accumulate CPU credits per hour, whose amount also depends on the instance type. If a burstable instance uses fewer CPU resources than required for baseline performance (for example, when it is idle), the unspent CPU credits are accrued in the CPU credit balance of the instance. If the latter needs to burst above the baseline performance level, it spends the accrued credits. The higher the number of credits that a burstable instance has accumulated, the longer it can burst beyond its baseline when higher performance is required. Burstable instances have two main advantages: i) they are offered with an up to 20% discount compared to non-burstable on-demand instances with equivalent computational resources, and ii) contrarily to spot VMs, they are not prone to revocation neither hibernation. On the other hand, to obtain monetary advantages of burstable instances, the user has to control their respective CPU credit usage by monitoring their baseline performance and defining bursting periods. In Amazon EC2, burstable VMs are of type T3, T3a, or T2 [1].

In the current work, we consider Bag-of-Tasks (BoT) applications, which are commonly used in solving various science and engineering problems, such as parameter sweep, chromosome mapping, Monte Carlo simulation and computer imaging applications [2]. As the interconnection network is not the bottleneck for the good performance of this type of applications, it is possible to speed up them by distributing their tasks on a large number of VMs in the cloud. In this context, an important issue is the definition of an efficient initial scheduling for the tasks which minimizes both the execution time and the financial cost, typically conflicting objectives. To this end, we propose to mathematically formulate it as a multi-objective integer programming problem and/or solve it by applying metaheuristics efficiently. Furthermore, for many of such applications, cost optimization algorithms should respect a given deadline, otherwise, a temporal failure takes place. Considering that the application executes on spot VMs that can be revoked/hibernated, we should provide a second scheduler, denoted dynamic scheduler, which would be responsible for provisioning new VMs and migrating tasks between different VMs, if necessary.

Therefore, this article presents a framework for scheduling Bag-of-Tasks (BoT) applications with deadline constraints, exploring hibernate-prone spot instances and on-demand burstable instances aiming at minimizing both the monetary cost and the execution time of applications. The proposed framework, denoted Burst Hibernation-Aware Dynamic Scheduler (Burst-HADS), is an event-driven scheduling framework built in a modular way with two main scheduling modules: i) the Primary Scheduling Module that defines an initial scheduling map of tasks to VMs, and ii) a Dynamic Scheduling Module responsible for task migration in case of hibernation.

In a previous work [3], we have explored the use of hibernation-prone spot instances to minimize the monetary

cost of BoT applications execution, respecting their deadline constraints. To meet such a deadline, even in the presence of multiple hibernations, new on-demand VMs, not allocated in the initial mapping, could be dynamically launched. In this case, tasks of the hibernated spot instances and those not executed yet are migrated to on-demand instances. Although the strategy can significantly reduce the monetary costs, always respecting the application deadline, if VM spots hibernate, the application's total execution time may considerably increase when compared to a solution that uses only on-demand instances. Hence to tackle such a problem, in this article, we investigate how burstable instances can be used to reduce the impact on the execution time of the application and the corresponding monetary costs. We also propose a formulation of the task scheduling problem as a multi-objective integer programming problem (execution time and financial cost objectives), and a heuristic based on the Iterated Local Search (ILS) [4] method for generating the initial scheduling map.

Applying different scenarios, we have studied and evaluated our proposed framework. The results show that the burstable instance based approach can optimize the monetary cost without degrading the application execution time, even in the presence of several hibernations. The evaluation of Burst-HADS targeted different issues. First, we focused on the new proposed Iterated Local Search (ILS) based primary scheduler. The ILS approach was compared to optimal solutions given by the exact model of the problem and also with three widely used scheduling heuristics: MinMin, MaxMin and Greedy [5]. Next, we conducted a series of experiments with a real scenario using the Amazon EC2 cloud [6], where Burst-HADS was compared to three other approaches: (i) a baseline case where only on-demand VMs are used, (ii) HADS, proposed in our previous work [3], and (iii) an approach based on AutoBoT [7], a BoT scheduler that also exploits spot VMs to minimize the monetary cost while meeting the deadline defined by the user. To the best of our knowledge, AutoBoT is the closest to our work in the literature. However, since it considers the old EC2 price model, which was based on price history and spot market bid, and its original code is not available, we have implemented a version of AutoBot, denoted AutoBot-like, adapting some procedures of it to the current EC2 price model, defined by the supply and demand of spare resources. Performance results with experiments using both real and synthetic BoT applications in different scenarios subject to several spot VM hibernations and resuming events confirm the effectiveness of Burst-HADS in terms of monetary costs and execution times when compared to the other approaches.

The remainder of this paper is organized as follows. Section 2 discusses some related work, including the description of AutoBot and HADS, used in our analysis. Section 3 presents the mathematical formulation and the primary and dynamic modules of the Burst-HADS framework. The evaluation results of both the ILS based primary scheduling module and the dynamic module are presented in Section 4. Finally, Section 5 concludes the paper and introduces some future directions.

2 RELATED WORK

Since AWS introduced the concept of burstable VM, many works exploring its features have been proposed. Many of them focus on evaluating the performance gain obtained by

using burstable VMs. In [8], Leitner and Scheuner presented a first empirical and analytical study about the second generation of AWS burstable instances (T2 family). They considered the T2.micro, T2.small, and T2.medium instances. Their work aimed at answering if, in terms of monetary cost and performance, these instance types are more effective than other ones. The results show that compared to general-purpose and computed-optimized instances (2015 generation), the evaluated T2 instances provide a higher CPU performance-cost ratio as long as the average utilization of instances is below 40%. To figure out the CPU usage limits on-the-fly, considering the dynamic variation of workloads, Ali *et al.* proposed in [9] an autonomic framework that combines light-weight profiling and an analytical model. The objective was to maximize the amount of work done using the burstable capacity of T2 VM instances. The authors state that the framework extends the CPU credits depletion period. Similarly to Leitner and Scheuner's work [8], their results also confirmed the benefits of active CPU usage control when burstable instances are exploited.

Jiang *et al.* [10] analytically modeled the performance of burstable VMs, considering their configurations, such as CPU, memory, and CPU credits parameters. They also showed that providers could maximize their total revenue by finding the optimal prices for burstable instances. Although their work, contrarily to ours, does not focus on application performance, its contribution is interesting since it states that providers can offer burstable instances for low prices without losing revenue while meeting QoS parameters.

Other works take advantage of burstable instance features in scheduling and scaling problems. In [11], for example, Baarzi *et al.* proposed an autoscale tool denoted BurScale, which uses burstable instances, together with on-demand ones, in order to handle transient queuing which arises due to traffic variation. They also presented how burstable instances can mask VM startup/warmup costs when autoscaling, which handles flash crowds, takes place. Using two distinct workloads, a stateless web server cluster and a stateful Memcached caching cluster, the authors showed that a careful combination of burstable and regular instances ensures similar performance for applications as traditional autoscaling systems while reducing up to 50% of the monetary cost. In [12], Wang *et al.* combined on-demand, spot, and burstable instances proposing an in-memory distributed storage approach. Burstable instances were used as a backup to overcome performance degradation resulting from spot instance revocations. According to the authors, those instances' burst capacity makes them ideal candidates for such a backup. Performance results show that the backup using burstable instances presents latency 25% lower than the latency of backup based on regular instances, inducing, therefore, a significant monetary cost saving.

Spot and on-demand instances have received a lot of attention in applications scheduling [7], [13], [14], [15], [16], [17]. Most of those works cope with termination/revocation of spot VMs and do not consider the hibernation feature. Moreover, the majority of them exploit the historical of spot VM price variation to predict spot VMs' revocations. However, since EC2 adopted a new price model (announced in December 2017), prices of VMs in the spot market are quite stable and defined exclusively by the supply and demand of spare capacity and no more by bid prices.² Therefore, it is

no longer possible to predict the revocation/termination of spot VMs based only on the history of price variations as those works do.

To the best of our knowledge, the hibernation mechanism of the spot VMs is only discussed in our previous works [3], [16] and in Fabra *et al.* [18]. We proposed, in [16], a static heuristic that creates predefined backup maps before the execution of the job tasks themselves. It was the first attempt to cope with spot VMs hibernation, and results from simulations showed that the hibernation problem is better handled with a dynamic approach. In [3], we present a dynamic scheduler, denoted HADS that uses both spot and regular (non-burstable) on-demand VMs to execute BoT applications. It aims at minimizing the execution's monetary cost respecting application deadline. In [18], the authors consider a scenario where hibernation-prone spot VMs can be used and then they show that deadline constraints add complexity to the problem of resource provisioning. However, no performance experiment results are presented neither discussed. Moreover, the work does not focus on the task scheduling problem, but on resource provisioning.

Recently, in [7], Varshney and Simmhan proposed the AutoBoT Scheduler. According to the authors, AutoBoT uses several heuristics to choose a set of spot and on-demand VMs to execute BoT applications subject to a deadline defined by the user. The objective is to reduce the monetary cost of the execution, and the scheduler combines checkpoints and migration procedures to meet the deadline, even with spot revocations. Like Burst-HADS, AutoBoT also has two execution phases: i) the initial mapping, where an initial execution map is built, and ii) the execution phase, where the tasks are executed. Similarly to the Dynamic Scheduling of Burst-HADS, the execution phase of AutoBoT also reacts to the spot revocations by changing the initial execution map, adding on-demand VMs, and migrating the affected tasks. However, there are some crucial differences between AutoBoT and Burst-HADS. First, although it is a recent work, AutoBoT takes some scheduling decisions based on the old spot price model. Moreover, differently from Burst-HADS, AutoBoT does not explore the hibernation feature of spot VMs neither the burst capacity of burstable VMs to reduce the monetary cost and minimize the execution time. Another important point to highlight is that AutoBoT was not evaluated in a real cloud. Instead, the authors presented a comprehensive study in a simulated environment.

3 SCHEDULING WITH HIBERNATION-PRONE SPOT AND BURSTABLE ON-DEMAND INSTANCES

The Burst Hibernation-Aware Dynamic Scheduler (Burst-HADS) has two main scheduling modules: i) the Primary Scheduling Module, that defines an initial scheduling map (Section 3.2); and ii) a Dynamic Scheduling Module responsible for migrating tasks or executing task rollback recovering in idle VMs (Section 3.3). In this section, we present the algorithms and procedures proposed for each one of these modules. In Section 3.1, we formulate the primary scheduling map as a multi-objective integer programming problem.

3.1 Mathematical Formulation

The problem of scheduling tasks in distributed computing resources is an NP-complete one [19], even in simple

2. <https://aws.amazon.com/pt/blogs/compute/new-amazon-ec2-spot-pricing/>

scenarios. Furthermore, some features of clouds render it more difficult. We thus model the primary task scheduling problem as a multi-objective integer programming problem whose objectives are to minimize both the monetary cost and the total execution time of the application. In other words, in our case, it is defined as the problem of creating an initial scheduling strategy, respecting the limit of available virtual cores (vCPUs) and memory capacity of the used VMs while minimizing the makespan and the monetary costs of the execution. Then, a number of burstable VMs, defined in function of the number of selected spot VMs, are included as an additional resource at the last step of the Primary Scheduling Module (see Section 3.2). Therefore, in the mathematical model we consider only spot VMs.

Let M^s be the set of spot VMs that can be used to execute the BoT application and let D be the application deadline, defined by the user. We then consider a set $T = \{1, \dots, D\}$ of discrete valid time periods. Each $vm_j \in M^s$ has a memory capacity of m_j gigabytes and nvc_j vCPUs. When a new vm_j is launched, the user is charged c_j for each period of time. When the VM terminates or hibernates, the user's charge for this VM immediately stops. Let B be the set of tasks of a BoT application. We assume that each task $t_i \in B$ requires a known amount of memory rm_i , and it is executed in only one vCPU of vm_j . A multi-core vm_j ($nvc_j > 1$) can execute more than one task simultaneously (one task per core) only when there is enough main memory to allocate them in the VM. We also consider that the execution time e_{ij} of each task t_i in any $vm_j \in M^s$ is known beforehand.

Aiming at ensuring the application deadline D no matter if and when spot VM hibernations happen, in our previous work [3], we define D_{spot} as the worst-case estimated makespan, which guarantees that there will always have enough spare time to migrate tasks of any hibernated spot VM to other VMs and to execute them before the deadline D . D_{spot} is computed by considering the deadline D and the longest tasks that might need to be migrated and executed to/in the slowest VMs. Let the binary variable X_{ij}^v indicate whether a task $t_i \in B$ allocated to a $vm_j \in M^s$ will start executing ($X_{ij}^v = 1$), or not ($X_{ij}^v = 0$), at time period $v \in T$. Let also Z_j and ZT be continuous variables which respectively keep the last period of execution of a $vm_j \in M^s$ and the total execution time of the application (makespan).

The proposed objective function (Equation (1)) is a weighted function that minimizes the monetary cost and the makespan, where α is the weight given by the user for the objectives.

$$\min \left(\alpha \times \left(\sum_{vm_j \in M^s} Z_j \times c_j \right) + (1 - \alpha) \times ZT \right). \quad (1)$$

Note that, in this case, both the monetary cost and the makespan have to be first normalized. The normalization procedure updates the target values to share the same minimum and maximum values, 0 and 1, respectively. Thus, the solution's total monetary cost was divided by the product of the monetary cost of hiring the most expensive spot $vm_j \in M^s$, during D_{spot} periods, times the maximum number of VMs that can be deployed. Similarly, the makespan is divided by D_{spot} . When the heuristic finds a solution with makespan longer than the D_{spot} limit, it does not accomplish the normalization, and that solution is

discarded by ILS. The objective function is subject to the following constraints.

Constraint (2) guarantees that every task $t_i \in B$ must be executed, starting at a time $v \in T$ in a $vm_j \in M^s$. Constraint (3) ensures that tasks' memory demand does not outpace the memory capacity of the VM, while Constraint (4) guarantees that the number of parallel tasks allocated to a $vm_j \in M^s$ does not exceed the number of virtual cores of the VM.

$$\sum_{vm_j \in M^s} \sum_{v \in T} X_{ij}^v = 1, \forall i \in B \quad (2)$$

$$\sum_{t_i \in B} \sum_{q=p}^v rm_i \times X_{ij}^q \leq m_j, \quad (3)$$

$$\forall vm_j \in M^s, \forall v \in T, \text{ and } p = \max(v - e_{ij}, 1)$$

$$\sum_{t_i \in B} \sum_{q=p}^v X_{ij}^q \leq nvc_j, \quad (4)$$

$$\forall vm_j \in M^s, \forall v \in T, \text{ and } p = \max(v - e_{ij}, 1).$$

Inequalities (5) and (6) relate the last period of execution of each $vm_j \in M^s$ with the application total execution time (makespan). Finally, Constraint (7) ensures that the application makespan does not exceed the D_{spot} value.

$$X_{ij}^v \times (v + e_{ij}) \leq Z_j \quad (5)$$

$$\forall t_i \in B, \forall vm_j \in M^s \text{ and } \forall v \in T$$

$$X_{ij}^v \times Z_j \leq ZT \quad (6)$$

$$\forall t_i \in B, \forall vm_j \in M^s \text{ and } \forall v \in T$$

$$Z_j \leq D_{spot}, \forall vm_j \in M^s. \quad (7)$$

As shown in Section 4, the execution time to solve the exact model is prohibitive even to small BoT applications. Hence, in order to find good solutions to the scheduling problem in an acceptable time, we adopt the heuristic approach presented in Section 3.2.

3.2 Primary Scheduling Module

In order to create the initial scheduling map according to the Mathematical Formulation presented in Section 3.1, we propose an Iterated Local Search (ILS) able to solve the problem in an acceptable time. The ILS [4] is a metaheuristic that aims at improving a final solution by sampling in a broad and distant neighborhood of candidate solutions and then applying a local search technique to refine solutions to their local optima. It explores a sequence of solutions created by perturbations of the current best solution to reach these distant neighborhoods.

Let M^b be the set of burstable on-demand VMs. After finding a scheduling map with the ILS, a second heuristic is applied to include burstable instances from M^b into the solution. In case of spot VM hibernations, the respective instances will be used, in burst mode, by the dynamic scheduler, as an attempt to minimize the impact of these hibernations in the monetary cost and/or the execution time. Therefore, the Primary Task Scheduling algorithm (Algorithm 1) has two parts: i) the Iterated Local Search, that solves the scheduling problem defined in Section 3.1 and ii) the burstable instances allocation, that includes

TABLE 1
Variables and Parameters Used in Algorithms 1, 2 and 3

Name	Description
B	Set of tasks
M^s	Set of spot VMs
M^b	Set of on-demand burstable VMs
D	Application deadline
D_{spot}	Estimated time limit which ensures that there will be enough spare time to migrate tasks
t_i	Task $t_i \in B$
vm_j	Virtual machines
$max_iteration$	Number of iterations of the ILS
max_failed	Tolerated number of iterations without improving the solution
$max_attempt$	Maximum number of attempts to find a better solution
$attempt$	Current attempt to find a better solution
RD_{spot}	Relaxed D_{spot} value
$relaxed_rate$	Percentage of increment of the RD_{spot} value
it	Current iteration of the ILS
it_{best}	Last iteration where a best solution was found
$swap_rate$	Percentage of tasks to be swapped
S	A scheduling solution
n	Number of burstable VMs added to the final solution
S_{final}	The final scheduling solution
$burst_rate$	Percentage of burstable VMs included into the final solution
A	Set of selected VMs

burstable instances to the final solution. Table 1 summarizes the variables and parameters used in the Primary Task Scheduling Algorithms 1, 2 and 3.

Initially, in line 2 of Algorithm 1, an initial solution is generated by calling Algorithm 2, originally proposed in [3]. The latter is a greedy heuristic that schedules the set of tasks B to a set M^s of spot VMs. First, in line 2 of Algorithm 2, all tasks of B are sorted in descending order by their memory size requirements. Then, by calling the procedure *check_schedule*, it verifies if it is possible to schedule task t_i in the selected vm_j of the current phase, i.e., the task is scheduled to the VM if the memory requirements are satisfied and if D_{spot} is not violated. Note that scheduling tasks in an already allocated VM avoids VM deploying time. Thus, for each task $t_i \in B$, the algorithm tries to schedule it to a vm_j from A , the set of already selected VMs (lines 7 to 12). If task t_i cannot be scheduled to a vm_j of A , the algorithm tries to select a new spot VM (lines 14 to 23) using a weighted round-robin algorithm (WRR) [20]. In WRR, each spot VM has an associated weight, and the algorithm selects the VMs in a round-robin way, according to such weights.

As shown in Equation (8), the weight of vm_j , $weight(vm_j)$, is equal to the quotient between $Gflops_j$ of vm_j , and c_j , the price of the VM per time period. In this work, the $Gflops_j$ of a vm_j is used to quantify the computing power of vm_j . It is previously estimated using the LINPACK benchmark [21] and allows us to compare the VMs' performance. Our choice in using WRR and spot VMs with different configurations is in agreement with Amazon's recommendations³ that say that an application should use

different types of spot VMs to increase the availability of spot VM instances. Interruptions of spot VMs, which include hibernation, usually take place in VMs of the same type. Therefore, a choice of heterogeneous spot VMs minimizes the impact of possible VM hibernations. Finally, in line 25, Algorithm 2 calls the function *create_solution* that receives the set of the selected VMs A and returns a solution S that defines a scheduling map.

$$weight(vm_j) = Gflops_j / c_j, \text{ where } vm_j \in M. \quad (8)$$

Algorithm 1. Primary Task Scheduling

Input: $B, M, M^s, M^b, max_iteration, max_attempt, max_failed, relaxed_rate, D_{spot}$ **and** D

```

1: { /*PART 01 - Iterated Local Search*/ }
2:  $S \leftarrow initial\_solution(B, M^s, D_{spot})$  {Algorithm 2}
3:    $S \leftarrow local\_search(S, max\_attempt, B, swap\_rate, D_{spot})$ 
   {Algorithm 3}
4:  $S_{best} \leftarrow S$ 
5:  $RD_{spot} \leftarrow D_{spot}$ 
6:  $it, it_{best} \leftarrow 0$ 
7: while  $it < max\_iteration$  do
8:    $vm_j \leftarrow random\_choice(M^s)$ 
9:    $S.selected\_vms \leftarrow S.selected\_vms \cup vm_j$ 
10:   $M^s \leftarrow M^s \setminus \{vm_j\}$ 
11:  if  $(it - it_{best}) > max\_failed$  then
12:     $RD_{spot} \leftarrow RD_{spot} + (relaxed\_rate \times RD_{spot})$ 
13:  end if
14:   $S \leftarrow local\_search(S, max\_attempt, B, swap\_rate, RD_{spot})$ 
15:  if  $fitness(S, RD_{spot}) < fitness(S_{best}, RD_{spot})$  then
16:     $S_{best} \leftarrow S$ 
17:     $it_{best} \leftarrow it$ 
18:  end if
19:   $it \leftarrow it + 1$ 
20: end while
21:
22: { /*PART 02: Burstable instance allocation*/ }
23:  $n \leftarrow \lceil burst\_rate \times S_{best}.selected\_vms \rceil$ 
24:  $S_{final} \leftarrow burst\_allocation(S_{best}, burst\_rate, M^b, D_{spot}, D)$ 
25:  $create\_primary\_map(S_{final})$ 
```

After obtaining an initial solution, the ILS algorithm tries to improve it by applying local search and perturbation procedures (Algorithm 1, lines 3 to 20). Thus, let S be a solution that defines a scheduling map of all tasks $t_i \in B$ to a subset of VMs of M^s . Let $fitness(S)$ be a weighted function that assigns a value to the quality of S . Since this function is equivalent to the objective function presented in Equation (1), we define in Equation (9) the $fitness(S, D_{spot})$ function, where $cost$ is the total monetary cost of S and mkp is the total execution time of the application.

$$fitness(S, D_{spot}) = \begin{cases} \infty, & \text{if violates } D_{spot} \\ \alpha \cdot cost + (1 - \alpha) \cdot mkp, & \text{otherwise} \end{cases} \quad (9)$$

Algorithm 1 executes a local search by calling, in line 3, the *local_search* procedure (Algorithm 3), which executes a series of attempts to improve the current solution by swapping tasks between the selected VMs. Algorithm 3 receives as input the current solution S , the *max_attempt* parameter that determines the number of times the local search will be

3. <https://aws.amazon.com/pt/ec2/spot/instance-advisor/>

executed, the set of tasks B , the $swap_rate$ and the D_{spot} value. The $swap_rate \in [0, 1]$ parameter is tuned before the execution and determines the number of tasks that will be swapped at each iteration. All the parameters used in our experiments, including the $swap_rate$, $max_attempt$, $max_iteration$ and other parameters will be presented in Section 4.

As can be observed in lines 4 and 8 of Algorithm 3, a solution S is composed of two structures: (i) a vector, which controls task allocation, where indexes correspond to tasks, and each element keeps the identity of the VM that will execute the corresponding task, and (ii) a list composed by selected VMs. First, the algorithm computes the number of tasks that will be swapped at each iteration (line 2) and randomly selects a destination VM (vm_{dest} , line 4). After that, the algorithm starts the tasks swapping procedure (lines 5 to 14), where n tasks, also randomly selected (line 7), are moved to the vm_{dest} . After each swap movement, the *local_search* procedure checks if the quality of the new generated solution has improved (line 9) and it updates the S_{best} solution, if necessary. In the end, the procedure returns the best solution (line 15).

Algorithm 2. Initial Solution

Input: B, M^s, D_{spot}

```

1:
2: sort_by_memory( $B$ ) {Sort Tasks  $t_i \in B$  according to  $rm_i$ }
3:  $A \leftarrow \emptyset$  {Set of selected VMs}
4: for all  $t_i \in B$  do
5:   {P1: Try to schedule the task in an already selected spot VM}
6:   sort_by_price( $A$ ) {Sort VMs  $v_j \in A$  according to  $c_j$ }
7:   for all  $vm_j \in A$  do
8:     if check_schedule( $t_i, vm_j, D_{spot}$ ) then
9:       schedule( $t_i, vm_j$ )
10:      break {/*Schedule next task*/}
11:     end if
12:   end for
13:   {P2: Try to schedule the task in a new spot VM}
14:   if not_scheduled then
15:     {Select a spot VM using the weighted round-robin heuristic}
16:      $vm_k \leftarrow \text{weighted\_RR}(M^s)$ 
17:     if check_schedule( $t_i, vm_k, D_{spot}$ ) then
18:       schedule( $t_i, vm_k$ )
19:        $A \leftarrow A \cup \{vm_k\}$  {Update the set of selected VMs}
20:        $M^s \leftarrow M^s \setminus \{vm_k\}$ 
21:       break {/*Schedule next task*/}
22:     end if
23:   end if
24: end for
25:  $S \leftarrow \text{create\_solution}(A)$ 
26: Return  $S$  {Returns a scheduling solution  $S$ }
```

After the first execution of the *local_search* procedure (Algorithm 1, line 3), Algorithm 1 has a loop that first performs a perturbation (lines 8 to 13) and then a new local search (line 14). The perturbation is responsible for diverting the metaheuristic from local optimal solutions. In the current work, we use two perturbation strategies. The first one includes a not selected spot $vm_j \in M^s$ into the current solution S (lines 8 to 10). The second one, called relaxing

perturbation, increases the D_{spot} bound value (lines 11 to 13). Note that the latter is executed only when the number of iterations without finding a better solution is higher than the *max_failed* parameter (line 11). In this case, the metaheuristic increases the D_{spot} limit in *relaxed_rate* percent, where *relaxed_rate* $\in [0, \dots, 1]$ is also a parameter defined by the user.

Upon finishing the ILS (Part 1), Algorithm 1 executes the *burst_allocation* procedure (Part 2, lines 22 to 24). In this procedure, n burstable VMs are included in the final solution, where n is a percentage, given by the parameter *burst_rate*, of the selected spot VMs of the best solution found by the ILS. For example, if 20 spot VMs were selected by the ILS, with *burst_rate* = 0.1, only 2 burstable VMs will be included.

Algorithm 3. Local Search

Input: $S, max_attempt, B, swap_rate$ and D_{spot}

```

1:  $S_{best} \leftarrow S$ 
2:  $n \leftarrow swap\_rate \times |B|$ 
3: attempt  $\leftarrow 0$ 
4:  $vm_{dest} \leftarrow \text{random\_choice}(S.selected\_vms)$ 
5: while attempt < max_attempt do
6:   for  $k \in \{1, \dots, n\}$  do
7:      $t_i \leftarrow \text{random\_choice}(B)$ 
8:      $S.allocation\_array[t_i] \leftarrow vm_{dest}$ 
9:     if fitness( $S, D_{spot}$ ) < fitness( $S_{best}, D_{spot}$ ) then
10:       $S_{best} \leftarrow S$ 
11:     end if
12:   end for
13:   attempt  $\leftarrow attempt + 1$ 
14: end while
15: return  $S_{best}$ 
```

Since the relaxed perturbation leads some tasks to violate the D_{spot} limit, the *burst_allocation* procedure also moves these tasks to the burstable instances. Each burstable VM can receive at most one task to be executed in baseline mode. However, if there still exist tasks violating D_{spot} and no available burstable VM, the procedure allocates them to the cheapest regular on-demand VMs. On the other hand, if a burstable VM remains idle, the task with the latest finishing time in the scheduling map is moved to it. Remark that our strategy of having a single task per burstable instance at a time, executing in baseline mode, induces CPU credits accumulation. Consequently, these burstable instances become the best candidates to receive tasks in case of spot VM hibernations.

3.3 Dynamic Scheduling Module

The Dynamic Scheduling Module is responsible for performing some actions in response to events, such as spot VM hibernation, resuming, and idleness, that may occur along with the execution.

Let BR , IR , HR , and TR be the set of *busy*, *idle*, *hibernated*, and *terminated* VMs respectively. Thus, Burst-HADS considers that a vm_j can be in one of the following states i) *busy*, if active and executing tasks ($vm_j \in BR$); ii) *idle*, if active but not executing any task ($vm_j \in IR$); iii) *hibernated*, if it has been hibernated by the cloud provider ($vm_j \in HR$); and iv) *terminated*, if the VM has terminated or it was not available at the beginning of the application execution ($vm_j \in TR$).

One of the main goals of the Dynamic Scheduling Module is to decide when a VM should terminate. On the one hand, as VMs are charged for each period of time (see Section 3.1), the user has an interest that a VM terminates as soon it becomes *idle*, reducing cost and avoiding extra ones. On the other hand, in some cases, it might be interesting to keep it because it can receive and execute tasks without incurring deploying overhead. To tackle such a trade-off, Burst-HADS has a termination policy where the allocation time is logically divided into units denoted Allocation Cycles (ACs). A vm_j that reaches the end of its current AC, denoted AC_{cur_j} , in *idle* state, is terminated.

Algorithm 4. Burst Migration Procedure

Input: Q_l , IR , BR , M^o , and D

```

1:  $Q_l \leftarrow \text{sort\_tasks}(Q_l)$  /* Prioritizes tasks with checkpoints */
2: for each  $t_i \in Q_l$  do
3:   {Attempt 1 - Try to migrate task to a Burstable IDLE VM}
4:   for each burstable  $vm_j \in IR$  do
5:      $rcc_{ij} \leftarrow \lceil e_{ij} / \text{burst\_period} \rceil$ 
6:     if  $cc_j > rcc_{ij}$  and  $\text{check\_migration}(t_i, vm_j, D)$  then
7:       {Migrate  $t_i$  to burstable  $vm_j$  on burst mode}
8:        $\text{migrate}(t_i, vm_j)$ 
9:        $IR \leftarrow IR \setminus \{vm_j\}$ 
10:       $BR \leftarrow BR \cup \{vm_j\}$ 
11:      break /*Migrate next task*/
12:     end if
13:   end for
14:   {Attempt 2 - Try to migrate task to a NON-burstable Idle or Busy VM}
15:   if not migrated then
16:     /* Prioritizes idle spot VMs */
17:      $K \leftarrow \text{sort\_by\_market}(IR \cup BR)$ 
18:     for each NON-burstable  $vm_j \in K$  do
19:       if  $\text{check\_migration}(t_i, vm_j, D)$  then
20:          $\text{migrate}(t_i, vm_j)$ 
21:         if  $vm_j \in IR$  then
22:            $IR \leftarrow IR \setminus \{vm_j\}$ 
23:            $BR \leftarrow BR \cup \{vm_j\}$ 
24:         end if
25:         break /*Migrate next task*/
26:       end if
27:     end for
28:   end if
29:   {Attempt 3 - Migrate task to a new NON-burstable on-demand VM}
30:   if not migrated then
31:      $\text{sort\_by\_price}(M^o)$ 
32:     for each  $vm_j \in M^o$  do
33:       if  $\text{start}_{ij} + e_{ij} + \omega < D$  then
34:          $\text{start\_vm}(vm_j)$ 
35:          $\text{migrate}(t_i, vm_j)$ 
36:          $M^o \leftarrow M^o \setminus \{vm_j\}$ 
37:          $BR \leftarrow BR \cup \{vm_j\}$ 
38:         break /*Migrate next task*/
39:       end if
40:     end for
41:   end if
42: end for
  
```

The events and respective actions handled by the Dynamic Scheduling Module are the following:

- *vm_j becomes idle*: When a vm_j finishes the execution of all tasks scheduled to it, its state is updated to *idle* and, if there exist scheduled tasks that have not completely executed, a work-stealing procedure assigning some or all of these tasks to vm_j , whose state is, thus, set to *busy* (see Section 3.5).
- *Non-burstable idle vm_j reaches the end of AC_{cur_j}* : If at the end of AC_{cur_j} a non-burstable vm_j remains *idle*, it is terminated and removed from the set IR of *idle* VMs and included into TR , the set of *terminated* VMs.
- *Spot vm_j hibernates*: The vm_j 's state is updated to *hibernated* and if it was *busy*, Burst-HADS starts the *burst migration procedure* (see Section 3.4).
- *Spot vm_j resumes*: When a hibernated spot vm_j resumes, it is excluded from HR , the set of hibernated VMs, and Burst-HADS starts a work-stealing procedure that tries to move tasks from *busy* VMs to vm_j (see Section 3.5).

Table 2 summarizes the variables and parameters used on Algorithms of the Dynamic Schedule Module (Algorithms 4 and 5).

3.4 Burst Migration Procedure

As previously explained, as soon as a spot vm_l hibernates, Burst-HADS executes the migration procedure, searching for a set of VMs to assign and execute non-finished tasks, that were previously scheduled to vm_l . These tasks are denoted affected tasks.

Algorithm 4 presents the migration procedure, which always respects the deadline D when selecting tasks to migrate. It receives as input the set $Q_l \subset B$ of affected tasks, the sets of *idle*, *busy*, and *non-launched regular on-demand* VMs (IR , BR and M^o , respectively), and the deadline D . Note that, in the case of a burstable $vm_j \in M^b$, e_{ij} is the execution time of task t_i in vm_j in burst mode (100% of the vm_j processing power) and each $vm_j \in M^b$ have a current CPU credit amount cc_j that is constantly updated by the cloud provider (in the case of $vm_j \notin M^b$, i.e., non-burstable VMs, $cc_j = \infty$).

Initially, Q_l is ordered, giving priority to those tasks that were executing at the moment of the hibernation and had been previously checkpointed (line 1). Burst-HADS applies the checkpoint/recovery mechanism that we proposed and evaluated in our previous work [22] where checkpoints are taken with the help of Checkpoint Restore In Userspace (CRIU) [23], a widely used tool that records the state of individual applications. In order to avoid the overhead of launching new VMs, the migration procedure gives priority to the use of already launched VMs. For each task $t_i \in Q_l$, the algorithm first tries to migrate the task to an *idle* burstable VM (lines 4 to 13). Otherwise, it tries to schedule t_i to one of the non-burstable *idle* or *busy* VM of set $K = IR \cup BR$ (lines 18 to 27). Note that, the set K is created by sorting the *idle* and *busy* VMs according to its market, putting the spot VMs in the front of the ordered set K (Algorithm 4 line 17).

Tasks migrated to burstable VMs are executed in burst mode, i.e., using 100% of the VM's CPU processing power. Therefore, it is necessary to ensure that a selected burstable VM will have enough CPU credits to execute all tasks

TABLE 2
Variables and Parameters Used on Algorithms 4 and 5

Name	Description
Q_l	Set of unfinished tasks of a hibernated spot vm_l
IR	Set of <i>idle</i> VMs
BR	Set of <i>busy</i> VMs
M^o	Set of on-demand VMs
D	Application deadline
cc_j	Current amount of CPU credit
rcc_{ij}	Estimated number of required CPU credits
$burst_period$	Number of periods of T corresponding to the use of one CPU credit
$K = IR \cup BR$	Ordered set of idle and busy VMs
$start_{ij}$	The time period a task t_i will start executing when allocated to a vm_j
e_{ij}	Time required to execute task t_i in a vm_j
ω	Time overhead to deploy a new VM
vm_j and vm_k	Virtual machines
ST_j	Set of tasks that can be stolen from vm_j

assigned to it. Let $burst_period$ be the number of periods of T corresponding to one credit consumption in burst mode. Since we consider that a task t_i is executed in only one core (see Section 3.1), it is possible to estimate (line 5) the number of CPU credits consumed by task t_i , where e_{ij} is the execution time of t_i on vm_j in burst mode, and rcc_{ij} is the estimated number of required CPU credits. Then, the algorithm checks if vm_j has enough CPU credits and call the *check_migration* function to guarantee that vm_j has enough memory and will execute t_i before the deadline D (line 6). If both conditions are satisfied, t_i is migrated to vm_j in burst mode (line 8). In this case the burstable vm_j is removed from set IR and included in set BR (lines 9 and 10). Note that the cloud provider continuously updates the number of credits, cc_j , of every burstable VM.

Similarly, when Algorithm 4 migrates a task t_i to a non-burstable *idle* or *busy* vm_j , it also calls the function *check_migration* to verify if the VM has enough memory and will be able to finish the task, respecting the deadline D (line 19). However, if vm_j is a spot VM, the function *check_migration* should also verify if there will be enough spare time in vm_j between the end of the execution of vm_j tasks (including task t_i) and the deadline D since, in this case, a *busy* or *idle* spot vm_j is also subject to hibernation. The spare time has to be greater than the execution time of the longest task scheduled to vm_j , ensuring, therefore, that if a hibernation occurs, there will be enough time to migrate and execute all affected tasks before the deadline D .

Finally, if there does not exist any available already deployed VM able to execute task t_i , the algorithm migrates the task to a new on-demand VM of set M^o (lines 32 to 40). In this case, it is necessary to verify that, considering the start period of t_i in vm_j ($start_{ij}$), plus its execution time (e_{ij}) will not violate the deadline (line 33). The new allocated on-demand VM is then removed from set M^o and included in set BR (lines 36 and 37).

3.5 Work-Stealing

The work-stealing procedure, presented in Algorithm 5, aims at reducing the allocation time duration of regular on-demand VMs as well as balancing the load of spot VMs. It is triggered

when a hibernated spot VM resumes or when a VM (spot or on-demand) becomes *idle*. Basically, the procedure tries to move tasks from non-burstable *busy* VMs to the *idle* VM.

Algorithm 5. Burst Work-Stealing Procedure

Input: BR, IR, vm_k, D

```

1:  sort_by_market( $BR$ )  /*Prioritizes non-burstable on-demand VMs*/
2:  for each NON-burstable  $vm_j \in BR$  do
3:     $ST_j \leftarrow selectTasks(vm_j)$  {Select the tasks that can be stolen/}
4:    for each  $t_i \in ST_j$  do
5:      if check_migration( $t_i, vm_k, D$ ) then
6:        if  $vm_k$  is burstable then
7:          set_baseline_mode( $t_i, vm_k$ )
8:          migrate( $t_i, vm_k$ )
9:          stops the loop
10:         else if  $vm_k$  not burstable then
11:           migrate( $t_i, vm_k$ )
12:         end if
13:       end if
14:     end for
15:  end for
16:  if at least one task was stolen then
17:     $BR \leftarrow BR \cup \{vm_k\}$ 
18:     $IR \leftarrow IR \setminus \{vm_k\}$ 
19:  end if
```

For each non-burstable *busy* $vm_j \in BR$ the procedure selects the tasks that can be stolen from it (line 3) and tries to migrate them to the *idle* vm_k (lines 4 to 14). Since regular on-demand VMs are more expensive than spot ones, the procedure considers first the tasks from the former (line 1).

Similarly to the migration procedure, for each selected task of a vm_j , the work-stealing procedure also verifies, by calling the function *check_migration*, if the task migration would result in deadline violation (line 5). Since tasks migrated to burstable VMs by the work-stealing procedure are executed in the baseline mode, after verifying if the *idle* vm_k is burstable (line 6), the algorithm sets up the execution to the baseline mode (line 7). In this work, to set up a burstable VM to the baseline mode means that the task cannot run using 100% of the CPU processing power, but uses only the baseline performance defined by the provider. In order to limit the use of the CPU, Burst-HADS exploits the *cpulimit* tool,⁴ an open-source program that limits the CPU usage of a process in GNU/Linux. Moreover, if the *idle* vm_k is burstable, only one task is moved to it, to avoid a queue of tasks to be executed in the baseline mode, which could increase the application total execution time. Since stolen tasks are executed in baseline mode on burstable VMs non consumed CPU credits consuming keep available if a spot VM hibernates. Thus, after moving one task to the burstable vm_k , the algorithm stops the loop (line 9). Finally, if at least one task is migrated to vm_k , its state changes to *busy* and, consequently, it is included into the set of *busy* VMs and removed from the *idle* set (lines 17 and 18).

4 EXPERIMENTS AND RESULTS

In this section, we first evaluate the ILS, part one of Algorithm 1, in terms of quality of solution and execution time,

4. <http://cpulimit.sourceforge.net/>

TABLE 3
VMs Attributes

Type	#VCPUs	Memory	Price per Hour (USD)		Baseline performance
			on-demand	spot	
C3.large	2	3.75 GB	0.105\$	0.0299\$	-
C4.large	2	3.75 GB	0.100\$	0.0366\$	-
C3.xlarge	4	7.50 GB	0.199\$	0.0634\$	-
T3.large	2	8 GB	0.0832\$	-	20%

by comparing it with the solutions given by the mathematical formulation presented in Section 3.2 (solved with Gurobi 9.1) and with three commonly used BoT scheduling heuristics [24], [25], [26]: MinMin, MaxMin and Greedy (Section 4.1). We also present an evaluation in AWS EC2 (Section 4.2), where the jobs were executed in a real cloud environment. In the experiments, we have considered two distinct workloads: i) synthetic BoT applications composed by tasks generated with the application template proposed in Alves *et al.* [27], and ii) the ED application available in the GRIDNBP 3.1 suite of NAS benchmark [28]. The tasks of the synthetic applications execute vector operations whose execution times depend on the size of the vectors.

According to AWS, only spot VMs of types C3, C4, C5, M4, M5, R3, and R4 with memory amount smaller than 100 GB are hibernation-prone. Therefore, in our experiments, we used spot VMs C3 and C4, which are compute-optimized and have high availability in the spot market. Moreover, we also use the T3.large instances that are the newest generation of the general proposed burstable instances of EC2. Table 3 shows the computational characteristics and the corresponding prices in on-demand and spot markets in November 2020.

For the execution of Algorithms 1 and 3, the following input parameters were used: $\alpha = 0.5$, $max_iteration = 200$; $max_attempt = 50$; $swap_rate = 0.10$; $max_failed = 20$; $relaxed_rate = 0.25$; $burst_rate = 0.2$. Moreover, we defined $AC = 900$ seconds. Except for $\alpha = 0.5$, which was used to give the same weights to both objectives, the parameters' values were defined by executing a set of empirical tests. As mentioned in Section 3.4, Burst-HADS uses a checkpoint approach proposed in [3] where the user defines the parameter ovh as the maximum percentage of time overhead that the checkpoint mechanism is allowed to add to the original execution time of a task. From this parameter, Burst-HADS defines the time interval between checkpoints. In this work, $ovh = 10\%$ for all experiments.

4.1 Evaluation of the ILS Primary Task Scheduling

To evaluate the quality of the solution given by the ILS executed in the first part of the primary scheduling heuristic (Algorithm 1), we compared it with the optimal solution given by the exact model and also with the solutions given by three widely used benchmark scheduling heuristics [5]: (i) MinMin, where the scheduling gives priority to VMs with the minimum earliest completion time; (ii) MaxMin, where the scheduling gives priority to VMs with the maximum earliest completion time; and (iii) Greedy, where the shortest execution time task is scheduled to the cheapest available VM.

For the comparison with the exact model, we create six small synthetic jobs (from 5 to 30 tasks), composed of tasks with execution times varying from 5 to 15 seconds. We have chosen this set with up to 30 tasks because the exact model took more than ten hours to solve the problem for jobs with

TABLE 4
Results of the ILS-Based Primary Scheduling Heuristic and the Exact Approach

# Tasks	ILS			Exact Approach		
	Makespan	Cost	Time	Makespan	Cost	Time
5	21	\$0.0005	0.0071	21	\$0.0005	49.04
10	34	\$0.0011	0.0306	34	\$0.0011	70.10
15	34	\$0.0020	0.0297	34	\$0.0020	106.03
20	34	\$0.0021	0.0558	34	\$0.0021	233.10
25	34	\$0.0027	0.0736	34	\$0.0027	751.72
30	34	\$0.0035	0.0922	34	\$0.0035	9346.21

TABLE 5
Results of the ILS-Based Primary Scheduler, MinMin, MaxMin and Greedy Heuristics

# Tasks	ILS		MinMin		MaxMin		Greedy	
	mkp	cost	mkp	cost	mkp	cost	mkp	cost
50	57	\$0.0060	47	\$0.0073	610	\$0.0055	597	\$0.0051
100	86	\$0.0112	71	\$0.0220	897	\$0.0107	843	\$0.0101
150	109	\$0.0169	99	\$0.0338	885	\$0.0157	823	\$0.0148
200	138	\$0.0235	123	\$0.0376	898	\$0.0220	819	\$0.0205
250	163	\$0.0282	133	\$0.0461	888	\$0.0261	793	\$0.0246
300	204	\$0.0340	168	\$0.0551	899	\$0.0323	775	\$0.0317

more than 30 tasks. In these experiments, the deadline was fixed in 15 minutes ($D = 900$) for all jobs, and we used as input the spot VMs presented in Table 3, except for the burstable VM, since we are evaluating just the ILS heuristic. The experiments were executed on a computer with a processor Intel Core i7-3770 CPU 3.40 GHz with 12 GB of memory and Ubuntu 18.04 was used. The mathematical formulation was solved by using the Gurobi Solver 9.0.⁵

Table 4 summarizes the results with both the exact approach and the ILS-based heuristic. The first column identifies the number of tasks. The three columns that follow correspond to the results achieved by ILS: makespan, monetary cost, and the execution time to obtain the solution. The next three columns present the same results for the exact approach. The values shown for the ILS are averages of three executions. The ILS standard deviations were zero for all jobs.

As we can observe in Table 4, ILS obtains the same solutions of the exact approach for these small jobs, taking significantly less time than the mathematical formulation. On average, the ILS takes 0.048s against 2101.43s of the exact approach. Although these results are very encouraging, the experiments were limited to jobs with few tasks (up to 30 tasks) due to the huge required time to obtain the exact solution for jobs with more tasks.

For the evaluation with the other heuristics, six synthetic jobs were created with sizes varying from 50 to 300 tasks. The tasks execution times varied from 15 to 35 seconds, and the same deadline of 15 minutes ($D = 900$) was applied. Table 5 presents the makespan and monetary cost obtained by the ILS, MinMin, MaxMin, and Greedy heuristics. As shown in Table 5, the ILS heuristic outperforms MinMin in terms of monetary cost but increases the makespan. On average, the ILS makespan is increased of 18.11%, while the average monetary cost is reduced by more than 38.00%.

5. <http://www.gurobi.com>

TABLE 6
Characteristics of the Jobs

job	# tasks	memory footprint			task runtime(s)		
		min	avg	max	min	avg	max
J60	60	2.85 MB	4.69 MB	12.20 MB	102	198	323
J80	80	2.91 MB	4.71 MB	13.19 MB	103	199	322
J100	100	2.81 MB	4.49 MB	10.86 MB	107	190	330
ED200	200	153.74 MB	168.68 MB	177.77 MB	161	211	354

Compared to both MaxMin and Greedy, the ILS presents an average reduction in the makespan by more than 80%, with an average increment in the monetary cost by only 6.78% when compared to MaxMin, and 13.14% when compared to the Greedy heuristic. Contrarily to the MinMin heuristic, both MaxMin and Greedy heuristics reduce the solution's monetary cost since they give priority to the cheapest VMs, which increases the makespan.

It is worth noticing that both objectives were evenly considered in the ILS solutions and that the average loss in one of the objectives was always smaller than the gain in the other. Those results confirm the effectiveness of the proposed ILS to the BoT primary scheduling problem.

4.2 Evaluation of Burst-HADS on AWS EC2

For the experiments, several synthetic tasks were created, each one with a memory footprint between 2.81 MB and 13.19 MB, resulting in execution times which vary from 102 to 330 seconds. Then, we conceived three synthetic BoT applications, J60, J80, and J100 by randomly selecting the tasks. Concerning the ED application, a real embarrassingly distributed application, the job ED200 comprises 200 tasks solving the largest problem size (class B), taking 211 seconds per task on average.

Table 6 shows the four jobs' features, including their respective number of tasks, memory footprint and tasks runtime. In our experiments, we considered the shortest deadline in which Burst-HADS can find a feasible solution for all evaluated jobs, $D = 2700$ seconds (45 minutes), respecting the constraints. For some jobs, it was a tight deadline, while for other ones it was not. This approach allowed us to observe the behavior of the framework in those different scenarios.

We have first evaluated Burst-HADS in a scenario without hibernation, comparing it with (1) the schedule given by the proposed ILS using only on-demand VMs; (2) the schedule given by HADS, and (3) the schedule given by AutoBoT-like. The aim of these experiments is to measure the impact of including burstable on-demand VMs into the scheduling procedure in both the monetary cost and execution time. Table 8 presents the average of three executions of the

synthetic jobs J60, J80, and J100, and the real application ED200, for each case.

AutoBoT-like is inspired by the work presented in [7] (see Section 2). AutoBoT-like uses the same heuristics, procedures and parameters defined in [7], but the price history and the spot market bid are not considered in that adapted version. In our implementation, we use the current price model of AWS, which is more stable.

Initially, AutoBoT-like, as in the original version, considers a critical time point beyond which it is not possible to complete the remaining tasks without violating the deadline constraint, in case of a spot VM revocation. Thus, when a task, running on a spot VM, reaches this point, it is immediately migrated to an on-demand VM. The number of launched on-demand VMs in this case is defined in accordance with a sweep heuristic, based on the following parameters: the number of available on-demand VMs, the deadline and the total number of tasks of the application.

In the second step, along with the execution of the application, AutoBoT-like, as in the original version, takes checkpoints and migrate tasks to on-demand VMs. Although the authors propose three different strategies for checkpointing, we adopted the optimistic one, where checkpoints are recorded just before the critical time point, because the other ones depended on the price history and they did not present significant advantages.

In comparison with HADS, Table 7 shows that Burst-HADS reduces the makespan in 44.37%, 42.09%, 28.82%, and 11.82%, for jobs J60, J80, J100, and ED200, respectively. However, the average monetary cost increases by 66.34%, 44.54%, 57.55%, and 33.71%, for the same comparison. It happens because Burst-HADS already starts by using some burstable on-demand VMs. Moreover, the ILS based primary scheduling uses more VMs to reduce the execution time, while in HADS, the initial scheduling aims at minimizing only the monetary cost.

As can also observe in the same table, in the executions without spot revocations, Burst-HADS presents an average reduction of 20.92% in the monetary cost and 31.11% in the makespan, when compared to AutoBoT-like. Such a result is expected since AutoBoT-like launches additional on-demand instances even without spot revocations when the BoT execution gets the critical time point.

On the other hand, compared to the ILS on-demand strategy, on average, Burst-HADS reduces the monetary cost by more than 52.00%, with an average increase of 15% in the makespan. The ILS on-demand strategy uses the scheduling plan given by the ILS proposed in Algorithm 1, which includes neither spot nor burstable VMs, but only regular on-demand VMs. The longer makespan is due to the execution of tasks in the baseline mode of burstable VMs, which does not occur in the ILS on-demand strategy.

TABLE 7
Cost and Makespan of Burst-HADS, HADS (without hibernation), AutoBoT-like (without spot interruptions) and ILS On-Demand Only

JOB	Burst-HADS w/o Hibernation		HADS w/o Hibernation		AutoBoT-like w/o Interruptions		ILS On-demand	
	cost	makespan	cost	makespan	cost	makespan	cost	makespan
J60	\$0.112	1,274	\$0.067	2,290	\$0.166	2,221	\$0.271	1,112
J80	\$0.151	1,329	\$0.104	2,295	\$0.199	2,266	\$0.312	1,190
J100	\$0.176	1,660	\$0.112	2,332	\$0.218	2,342	\$0.371	1,462
ED200	\$0.357	2,275	\$0.267	2,580	\$0.387	2,566	\$0.698	1,887

TABLE 8
Different Execution Scenarios Generated by Varying
Parameters λ_h and λ_r

ID	hibernating	resuming	λ_h	λ_r
sc_1	$k_h = 1$	$k_r = 0$	1/2700	0/2700
sc_2	$k_h = 5$	$k_r = 0$	5/2700	0/2700
sc_3	$k_h = 1$	$k_r = 5$	1/2700	5/2700
sc_4	$k_h = 5$	$k_r = 5$	5/2700	5/2700
sc_5	$k_h = 3$	$k_r = 2.5$	3/2700	2.5/2700

Since cloud users have no control over spot VM hibernations, we have emulated different patterns of the events of spot hibernations and resume by using a Poisson distribution function. We have applied the Poisson function to modeling several scenarios where *hibernating* and *resuming* events have different probability mass functions defined by the parameters λ_h and λ_r , respectively. Let the λ parameter of Poisson distribution be the number of expected events divided by a time interval. Since the application execution is discretized by time interval and D is the application deadline, if we respectively define k_h and k_r , as the expected number (rate) of hibernating and resuming events during the application execution, λ_h and λ_r parameters are given by $\lambda_h = k_h/D$ and $\lambda_r = k_r/D$. Table 8 presents five different scenarios by varying k_h and k_r . Note that, since AutoBoT-like does not explore the hibernation feature of the spot VMs, instead of emulating the hibernation, the spot instances are terminated according to the Poisson function with probability mass functions also defined by the parameters λ_h .

Table 9 presents the averages of three executions of jobs J60, J80, J100, and ED200 using Burst-HADS, HADS and AutoBoT-like in each of the five execution scenarios. For each job and scenario, the table shows the average number of hibernations/revocations followed by resume events. It also includes the number of non-burstable on-demand VMs launched during the frameworks' execution and the average of monetary cost and makespan. Finally, the last four columns present the percentage differences related to the monetary cost and the makespan between Burst-HADS and HADS (Diff HADS), and between Burst-HADS and AutoBoT-like (Diff AutoBoT-like).

As shown in Table 9, in comparison to HADS, Burst-HADS minimizes the makespan in all execution scenarios, presenting an average reduction of 25.87%. As explained in Section 3.4, whenever a spot VM hibernates, Burst-

HADS immediately migrates the interrupted tasks to other VMs. Thus, the increase in the makespan is due to the overhead of this procedure, which might include the launch of new VMs. However, in these scenarios, the overhead has a lower impact in small jobs, i.e., jobs with fewer tasks, than in the ones with a higher number of tasks. For example, while for job J60, the average makespan reduction, considering all scenarios, is 40.10%, for job ED200, that reduction is only 10.24%. Such a behavior can be explained because, in our experiments, we have fixed the same deadline for all jobs and, therefore, small jobs have more spare time between its expected makespan defined by the ILS and the deadline. Consequently, in this case, Burst-HADS benefits more from the burst mode of the burstable VMs since it has more idle time to earn CPU credits. Moreover, it also executes the work-stealing more frequently, which also reduces the makespan. On the other hand, independently of the scenario or job, HADS's makespan gets closer to the deadline whenever a hibernation occurs. That happens because the HADS framework postpones as much as possible the execution of the migration procedure. Since HADS gives priority to the monetary cost saving, its central idea is to wait as much as possible for the resume of hibernated VMs in order to avoid the launch of new VMs.

Still compared to HADS, Burst-HADS improves the monetary cost for all jobs in scenarios sc_2 and sc_5 . The sc_2 is the worst case scenario, since it has the highest rate of hibernation ($k_h = 5$) and no resume ($k_r = 0$), while sc_5 is the average case scenario where the rate of hibernation is $k_h = 3.0$, and the rate of resume is $k_h = 2.5$ (see Table 8). In these scenarios, Burst-HADS uses fewer regular on-demand VMs than HADS. Moreover, in both cases, the number of hibernations is higher than the number of resumes. Hence, in those cases, by migrating tasks to busy and idle VMs and burstable VMs, exploring the burst mode, Burst-HADS is more effective in minimizing the impact of spot hibernations than HADS. It is also worth pointing out that, considering all executions, the average increase of Burst-HADS monetary cost is 1.92% compared to HADS.

Compared to AutoBoT-like, Burst-HADS can reduce the monetary cost for all jobs in almost all scenarios, presenting an average reduction of 11.65%. In our evaluation, AutoBoT-like reduced the monetary cost only in scenario sc_2 and for the jobs J60 and ED200. As explained before, this

TABLE 9
Comparison Between Burst-HADS, HADS and AutoBoT-Like in Terms of Monetary Cost and Makespan in Scenarios sc_1 to sc_5

Job	scenario	# hibernations	# resumes	# used regular on-demand VMs			Burst-HADS		HADS		AutoBoT-like		Diff HADS (%)		Diff AutoBoT-like (%)	
				Burst-HADS	HADS	AutoBoT-like	cost	makespan	cost	makespan	cost	makespan	cost	makespan	cost	makespan
J60	sc_1	0.66	0.00	0.00	0.00	1	\$0.119	1274	\$0.091	2620	\$0.166	2221	-30.77%	51.37%	28.31%	42.64%
	sc_2	3.33	0.00	1.33	2.33	1	\$0.204	1277	\$0.257	2549	\$0.173	2228	20.54%	49.90%	-17.92%	42.68%
	sc_3	2.33	2.33	1.33	0.00	1	\$0.127	1752	\$0.101	2539	\$0.178	2237	-26.07%	31.00%	28.46%	21.68%
	sc_4	5.33	4.00	1.67	0.00	1	\$0.142	1857	\$0.119	2634	\$0.180	2252	-19.90%	29.50%	21.07%	17.54%
	sc_5	2.66	1.00	1.33	2.00	1	\$0.150	1445	\$0.169	2359	\$0.170	2236	11.44%	38.75%	11.96%	35.38%
J80	sc_1	1.00	0.00	1.33	0.33	2	\$0.167	1419	\$0.150	2581	\$0.206	2273	-11.33%	45.03%	18.93%	37.57%
	sc_2	5.00	0.00	1.00	3.00	2	\$0.210	2267	\$0.298	2591	\$0.211	2278	29.48%	12.50%	0.47%	0.48%
	sc_3	3.00	1.00	1.67	1.00	2	\$0.164	1367	\$0.147	2602	\$0.214	2276	-11.34%	47.46%	23.52%	39.94%
	sc_4	9.66	7.66	1.00	2.00	3	\$0.244	2488	\$0.212	2607	\$0.301	2460	-15.25%	4.56%	18.83%	-1.14%
	sc_5	3.00	1.00	1.33	3.00	2	\$0.195	1589	\$0.246	2529	\$0.228	2266	20.47%	37.17%	14.25%	29.88%
J100	sc_1	2.00	0.00	0.00	0.00	2	\$0.191	1798	\$0.157	2332	\$0.226	2350	-21.76%	22.90%	15.49%	23.49%
	sc_2	7.00	0.00	1.33	3.00	2	\$0.212	1900	\$0.353	2518	\$0.222	2342	39.94%	24.54%	4.50%	18.87%
	sc_3	6.00	3.00	1.67	1.00	2	\$0.201	1925	\$0.166	2636	\$0.224	2360	-21.08%	26.97%	10.27%	18.43%
	sc_4	11.00	9.00	1.00	0.00	3	\$0.286	2453	\$0.278	2591	\$0.312	2677	-2.88%	5.33%	8.33%	8.37%
	sc_5	3.66	2.00	1.00	2.50	2	\$0.166	1547	\$0.189	2543	\$0.227	2366	12.49%	39.15%	26.98%	34.62%
ED200	sc_1	3.00	0.00	1.00	0.33	3	\$0.388	2327	\$0.314	2680	\$0.414	2630	-23.57%	13.17%	6.28%	11.52%
	sc_2	8.00	0.00	2.00	5.00	3	\$0.482	2448	\$0.512	2676	\$0.430	2661	5.86%	8.52%	-12.09%	8.00%
	sc_3	6.66	4.00	2.33	1.00	3	\$0.427	2345	\$0.387	2672	\$0.457	2675	-10.34%	12.24%	6.56%	12.34%
	sc_4	9.00	6.00	2.00	1.00	3	\$0.411	2560	\$0.389	2690	\$0.447	2667	-5.66%	4.83%	8.05%	4.01%
	sc_5	4.33	2.33	1.67	3.00	3	\$0.367	2342	\$0.467	2674	\$0.411	2765	21.41%	12.42%	10.71%	15.30%

TABLE 10

Comparison of Burst-HADS and AutoBoT-Like With an Unexpected Increment of 20% in the Execution Time of the Tasks of J60

Fluctuation (20%)				w/o Fluctuation			
Burst-HADS		AutoBoT-like		Burst-HADS		AutoBoT-like	
cost	makespan	cost	makespan	cost	makespan	cost	makespan
\$0.199	1,429	\$0.269	2,408	\$0.112	1,274	\$0.166	2,221
\$0.231	1,493	\$0.274	2,424	\$0.119	1,274	\$0.166	2,221
\$0.254	2,595	\$0.295	2,453	\$0.204	1,277	\$0.173	2,228
\$0.227	1,834	\$0.282	2,457	\$0.127	1,752	\$0.178	2,237
\$0.225	1,833	\$0.276	2,439	\$0.142	1,857	\$0.180	2,252
\$0.195	1,648	\$0.279	2,483	\$0.150	1,445	\$0.170	2,236

scenario has the highest rate of hibernation ($k_h = 5$), but there are no resume events ($k_r = 0$). Therefore, in this case, HADS and Burst-HADS cannot explore the resumed spot VMs to minimize the monetary cost. Although Burst-HADS is more effective than HADS to minimize that impact by using burstable VMs, for these two jobs, AutoBoT-like migrated the majority of tasks just before the VMs interruptions and selected a cheap set of on-demand VMs able to execute the remaining tasks. However, we should point out that the AutoBot-like approach is unsuitable for the other scenarios since it does not take advantage of the hibernation feature. In terms of makespan, Burst-HADS presented a better result in almost all cases when compared to AutoBoT-like, with an average reduction of 21.08%. The only exception for the makespan was scenario sc_4 for job J80, where there were many hibernations and resumes, which increased the overhead of Burst-HADS due to migrations and task work-stealing. But even in this case, the makespan difference was only 1.14%.

Finally, compared to the ILS On-demand, both Burst-HADS and HADS minimize the monetary cost for all execution cases, presenting an average reduction of 41.80% and 39.65%, respectively. For job ED200, for example, in the worst execution scenario, sc_2 , Burst-HADS reduces the monetary cost by 30.96%, while HADS presents an average reduction of 26.66%.

Since cloud environments can face random fluctuation in terms of performance, we also carried out some experiments considering an increase of 20% in the execution time of the tasks of job J60. We observed in Table 10 that both frameworks, AutoBoT-like and Burst-HADS, were able to meet the deadlines for all scenarios. However, just Burst-HADS kept the monetary cost below the on-demand only solution (\$0.271). However, for tight deadlines and depending on the range of the fluctuations, the frameworks might not meet the deadlines.

5 CONCLUSION AND FUTURE WORK

This paper proposes the Burst Hibernation-Aware Dynamic Scheduler (Burst-HADS) framework which uses regular on-demand VMs, hibernation-prone spot VMs, and burstable VMs for executing deadline constrained bag-of-task applications. The framework creates a primary scheduling plan using an ILS based heuristic and then executes a Dynamic Scheduling Module that monitors the execution, applying both task migration and work-stealing techniques whenever necessary. The objective of Burst-HADS is to minimize both the makespan and the monetary cost of the execution,

respecting applications deadline even when spot VMs in use hibernate.

Burst-HADS was evaluated in a real environment using VMs of AWS EC2, considering scenarios with different hibernation and resuming rates. When compared to the ILS On-demand approach that uses only regular on-demand VMs, Burst-HADS reduces the monetary cost for all execution scenarios at the expense of slightly longer makespans due to the migration overhead. Moreover, compared to HADS, our previous proposed framework that uses only spot and on-demand VMs, Burst-HADS reduces the makespan by more than 25%, with an average increase of only 1.92% in the monetary cost. Burst-HADS was also compared to AutoBoT-like, an approach inspired by Vashney and Simmhan's work. AutoBoT-like also combines spot and on-demand VMs but does not explore the hibernation feature. Our results show that, in this case, Burst-HADS reduces the average monetary cost and the makespan by 11.65% and 21.08%, respectively.

As future work, we would like to include in the migration decision of Burst-HADS some prediction of the future state of the spot market, considering, for example, the historic of VMs hibernation. A second research direction would focus on the selection of burstable VMs, with a deeper analysis of the most suitable one, based on the tradeoff between price and performance, in different execution scenarios.

REFERENCES

- [1] A. W. Services, "Burstable performance instances," Accessed: Jun. 2021. [Online]. Available: <https://docs.aws.amazon.com/burstable-performance-instances.html>.
- [2] L. Thai, B. Varghese, and A. Barker, "A survey and taxonomy of resource optimisation for executing bag-of-task applications on public clouds," *Future Gener. Comput. Syst.*, vol. 82, pp. 1–11, 2018.
- [3] L. Teylo, L. Arantes, P. Sens, and L. M. Drummond, "A dynamic task scheduler tolerant to multiple hibernations in cloud environments," *Cluster Comput.*, vol. 24, pp. 1051–1073, 2021.
- [4] H. R. Lourenço, O. C. Martin, and T. Stützle, "Iterated local search," in *Handbook of Metaheuristics*. Berlin, Germany: Springer, 2003, pp. 320–353.
- [5] J. O. Gutierrez-Garcia and K. M. Sim, "A family of heuristics for agent-based elastic cloud bag-of-tasks concurrent scheduling," *Future Gener. Comput. Syst.*, vol. 29, no. 7, pp. 1682–1699, 2013.
- [6] A. W. Services, "Amazon EC2 instance types," 2018. [Online]. Available: <https://aws.amazon.com/ec2/instance-types/>
- [7] P. Vashney and Y. Simmhan, "AutoBoT: Resilient and cost-effective scheduling of a bag of tasks on spot VMs," *IEEE Trans. Parallel Distrib. Syst.*, vol. 30, no. 7, pp. 1512–1527, Jul. 2019.
- [8] P. Leitner and J. Scheuner, "Bursting with possibilities—An empirical study of credit-based bursting cloud instance types," in *Proc. IEEE/ACM 8th Int. Conf. Utility Cloud Comput.*, 2015, pp. 227–236.
- [9] A. Ali, R. Pincioli, F. Yan, and E. Smirni, "CEDULE: A scheduling framework for burstable performance in cloud computing," in *Proc. IEEE Int. Conf. Auton. Comput.*, 2018, pp. 141–150.
- [10] Y. Jiang, M. Shahradd, D. Wentzlaff, D. H. Tsang, and C. Joe-Wong, "Burstable instances for clouds: Performance modeling, equilibrium analysis, and revenue maximization," in *Proc. IEEE Conf. Comput. Commun.*, 2019, pp. 1576–1584.
- [11] A. F. Baarzi, T. Zhu, and B. Urgaonkar, "BurScale: Using burstable instances for cost-effective autoscaling in the public cloud," in *Proc. ACM Symp. Cloud Comput.*, 2019, pp. 126–138.
- [12] C. Wang, B. Urgaonkar, A. Gupta, G. Kesidis, and Q. Liang, "Exploiting spot and burstable instances for improving the cost-efficacy of in-memory caches on the public cloud," in *Proc. 12th Eur. Conf. Comput. Syst.*, 2017, pp. 620–634.
- [13] S. Subramanya, T. Guo, P. Sharma, D. E. Irwin, and P. J. Shenoy, "SpotOn: A batch computing service for the spot market," in *Proc. 6th ACM Symp. Cloud Comput.*, 2015, pp. 329–341.
- [14] I. Menache, O. Shamir, and N. Jain, "On-demand, spot, or both: Dynamic resource allocation for executing batch jobs in the cloud," in *Proc. 11th Int. Conf. Auton. Comput.*, 2014, pp. 177–187.

- [15] P. Sharma, S. Lee, T. Guo, D. E. Irwin, and P. J. Shenoy, "SpotCheck: Designing a derivative IaaS cloud on the spot market," in *Proc. 10th Eur. Conf. Comput. Syst.*, 2015, pp. 16:1–16:15.
- [16] L. Teylo, L. Arantes, P. Sens, and L. M. d. A. Drummond, "A bag-of-tasks scheduler tolerant to temporal failures in clouds," in *Proc. 31st Int. Symp. Comput. Archit. High Perform. Comput.*, 2019, pp. 144–151.
- [17] T.-P. Pham and T. Fahringer, "Evolutionary multi-objective workflow scheduling for volatile resources in the cloud," *IEEE Trans. Cloud Comput.*, 2020, doi: [10.1109/TCC.2020.2993250](https://doi.org/10.1109/TCC.2020.2993250).
- [18] J. Fabra, J. Ezpeleta, and P. Alvarez, "Reducing the price of resource provisioning using EC2 spot instances with prediction models," *Future Gener. Comput. Syst.*, vol. 96, pp. 348–367, 2019.
- [19] J. D. Ullman, "NP-complete scheduling problems," *J. Comput. Syst. Sci.*, vol. 10, no. 3, pp. 384–393, 1975.
- [20] M. Katevenis, S. Sidiropoulos, and C. Courcoubetis, "Weighted round-robin cell multiplexing in a general-purpose ATM switch chip," *IEEE J. Sel. Areas Commun.*, vol. 9, no. 8, pp. 1265–1279, Oct. 1991.
- [21] J. J. Dongarra, P. Luszczek, and A. Petitet, "The linpack benchmark: Past, present and future," *Concurrency Comput.: Pract. Experience*, vol. 15, no. 9, pp. 803–820, 2003.
- [22] L. Teylo, R. C. Brum, L. Arantes, P. Sens, and L. M. d. A. Drummond, "Developing checkpointing and recovery procedures with the storage services of amazon web services," in *Proc. 49th Int. Conf. Parallel Process. Workshops*, 2020, pp. 1–8.
- [23] P. Emelyanov, "CRIU: Checkpoint/restore in userspace, July 2011," 2011. [Online]. Available: <https://criu.org>
- [24] N. Fujimoto and K. Hagihara, "A comparison among grid scheduling algorithms for independent coarse-grained tasks," in *Proc. Int. Symp. Appl. Internet Workshops*, 2004, pp. 674–680.
- [25] K. Li, "Job scheduling and processor allocation for grid computing on metacomputers," *J. Parallel Distrib. Comput.*, vol. 65, no. 11, pp. 1406–1418, 2005.
- [26] M. Maheswaran, S. Ali, H. J. Siegel, D. Hensgen, and R. F. Freund, "Dynamic mapping of a class of independent tasks onto heterogeneous computing systems," *J. Parallel Distrib. Comput.*, vol. 59, no. 2, pp. 107–131, 1999.
- [27] M. M. Alves and L. M. de Assumpção Drummond, "A multivariate and quantitative model for predicting cross-application interference in virtual environments," *J. Syst. Softw.*, vol. 128, pp. 150–163, 2017.
- [28] D. Bailey, T. Harris, W. Saphir, R. Van Der Wijngaart, A. Woo, and M. Yarrow, "The NAS parallel benchmarks 2.0," NASA Ames Research Center, Mountain View, CA, Tech. Rep. NAS-95-020, 1995.



Luan Teylo received the graduate degree in computer science, in 2015, the master's degree, in 2017, and the PhD degree in computer science from Fluminense Federal University, Brazil, in 2021. In 2019 he did a doctorate sandwich in the LiP6 Lab, Sorbonne Université, France, where he worked with scheduling and fault tolerance on the cloud. His research interests include I/O, distributed algorithms, cloud computing, and scheduling problems.



Luciana Arantes received the PhD degree in computer science from Paris 6 University (UPMC), France, in 2000. She has been working as an associate professor at Sorbonne University (ex-UPMC), France, since 2001. She was a research member of INRIA/LIP6 Regal project-team from 2005–2017 and now research member of INRIA/LIP6 Delys Group. She was member of the program committee of several conferences in the area of distributed systems and parallelism (ICDCS, EDCC, ISSRE, NCA, SBAC-PAD, DAIS, etc.) and local organized chair of SBAC-PAD 2014 and EDCC 2015.



Pierre Sens received the PhD degree in computer science, in 1994, and the habilitation *diriger des recherches* from Paris 6 University, France, in 2000. Currently, he is a full professor at Sorbonne Université, France. He was a member of the program committee of major conferences in the areas of distributed systems and parallelism (DISC, ICDCS, IPDPS, OPODIS, ICPP, Europar), and serves as general chair of SBAC and EDCC. Overall, he has published more than 150 papers in international journals and conferences and has acted for advisor of 24 PhD thesis.



Lúcia M. A. Drummond received the DSc degree in systems engineering and computer science from the Federal University of Rio de Janeiro, Brazil, in 1994, where she took part of the group which developed the first Brazilian parallel computer. She is in the Department of Computer Science of the Fluminense Federal University (UFF), Brazil, since 1989, where she is now full professor. She was member of the program committee of major conferences in the areas of high-performance computing and parallelism (ICS, CLUSTER, CCGRID, Europar, SBAC,) and served as program chair of SBAC twice, and guest editor of JPDC.

► For more information on this or any other computing topic, please visit our Digital Library at www.computer.org/csdl.